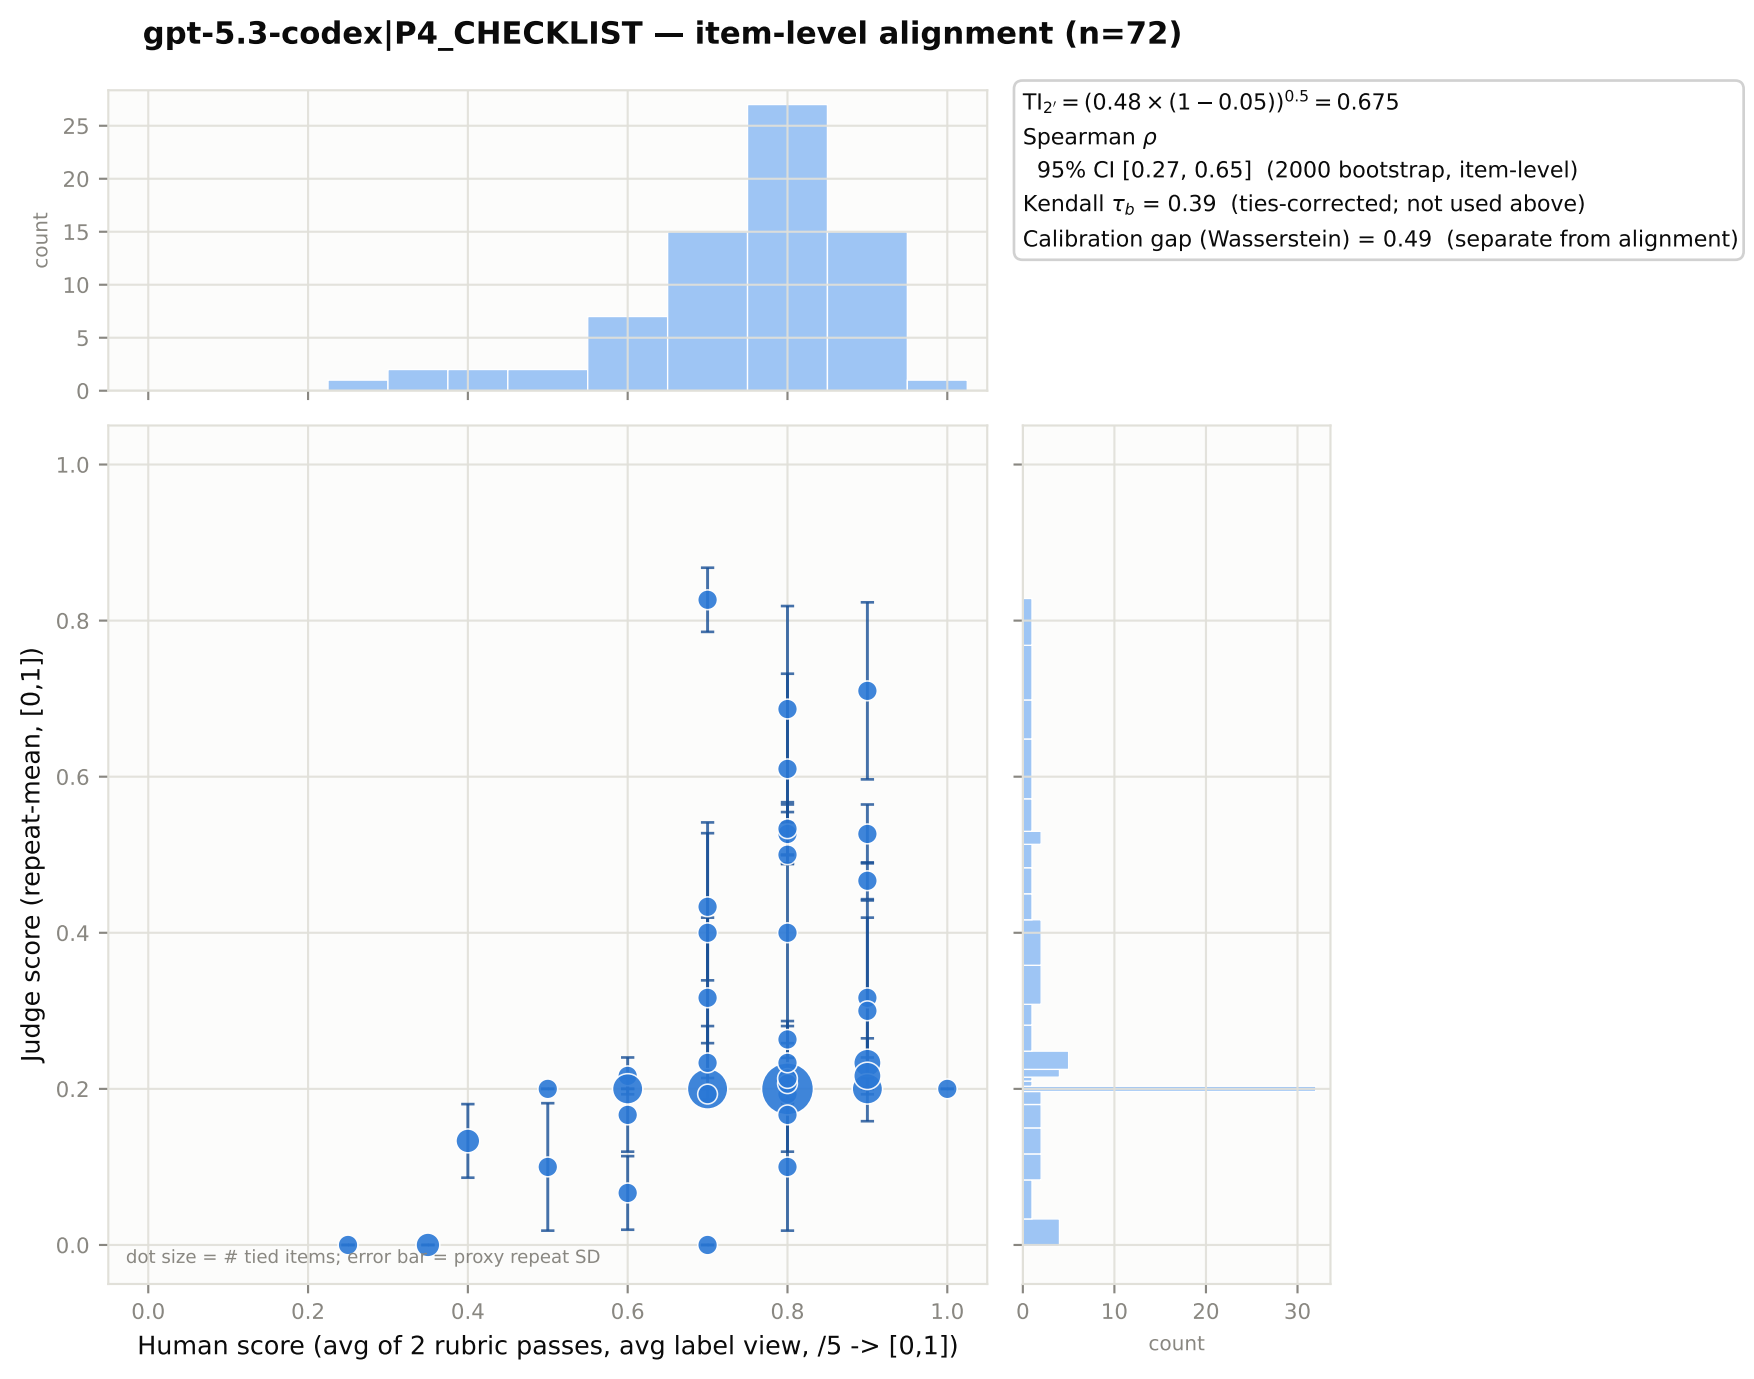
gpt-5.3-codex|P4_CHECKLIST — item-level alignment (n=72)
$TI_{2'} = (0.48 \times (1 - 0.05))^{0.5} = 0.675$
Spearman ρ
95% CI [0.27, 0.65] (2000 bootstrap, item-level)
Kendall τ_b = 0.39 (ties-corrected; not used above)
Calibration gap (Wasserstein) = 0.49 (separate from alignment)
count
Judge score (repeat-mean, [0,1])
dot size = # tied items; error bar = proxy repeat SD
Human score (avg of 2 rubric passes, avg label view, /5 -> [0,1])
count