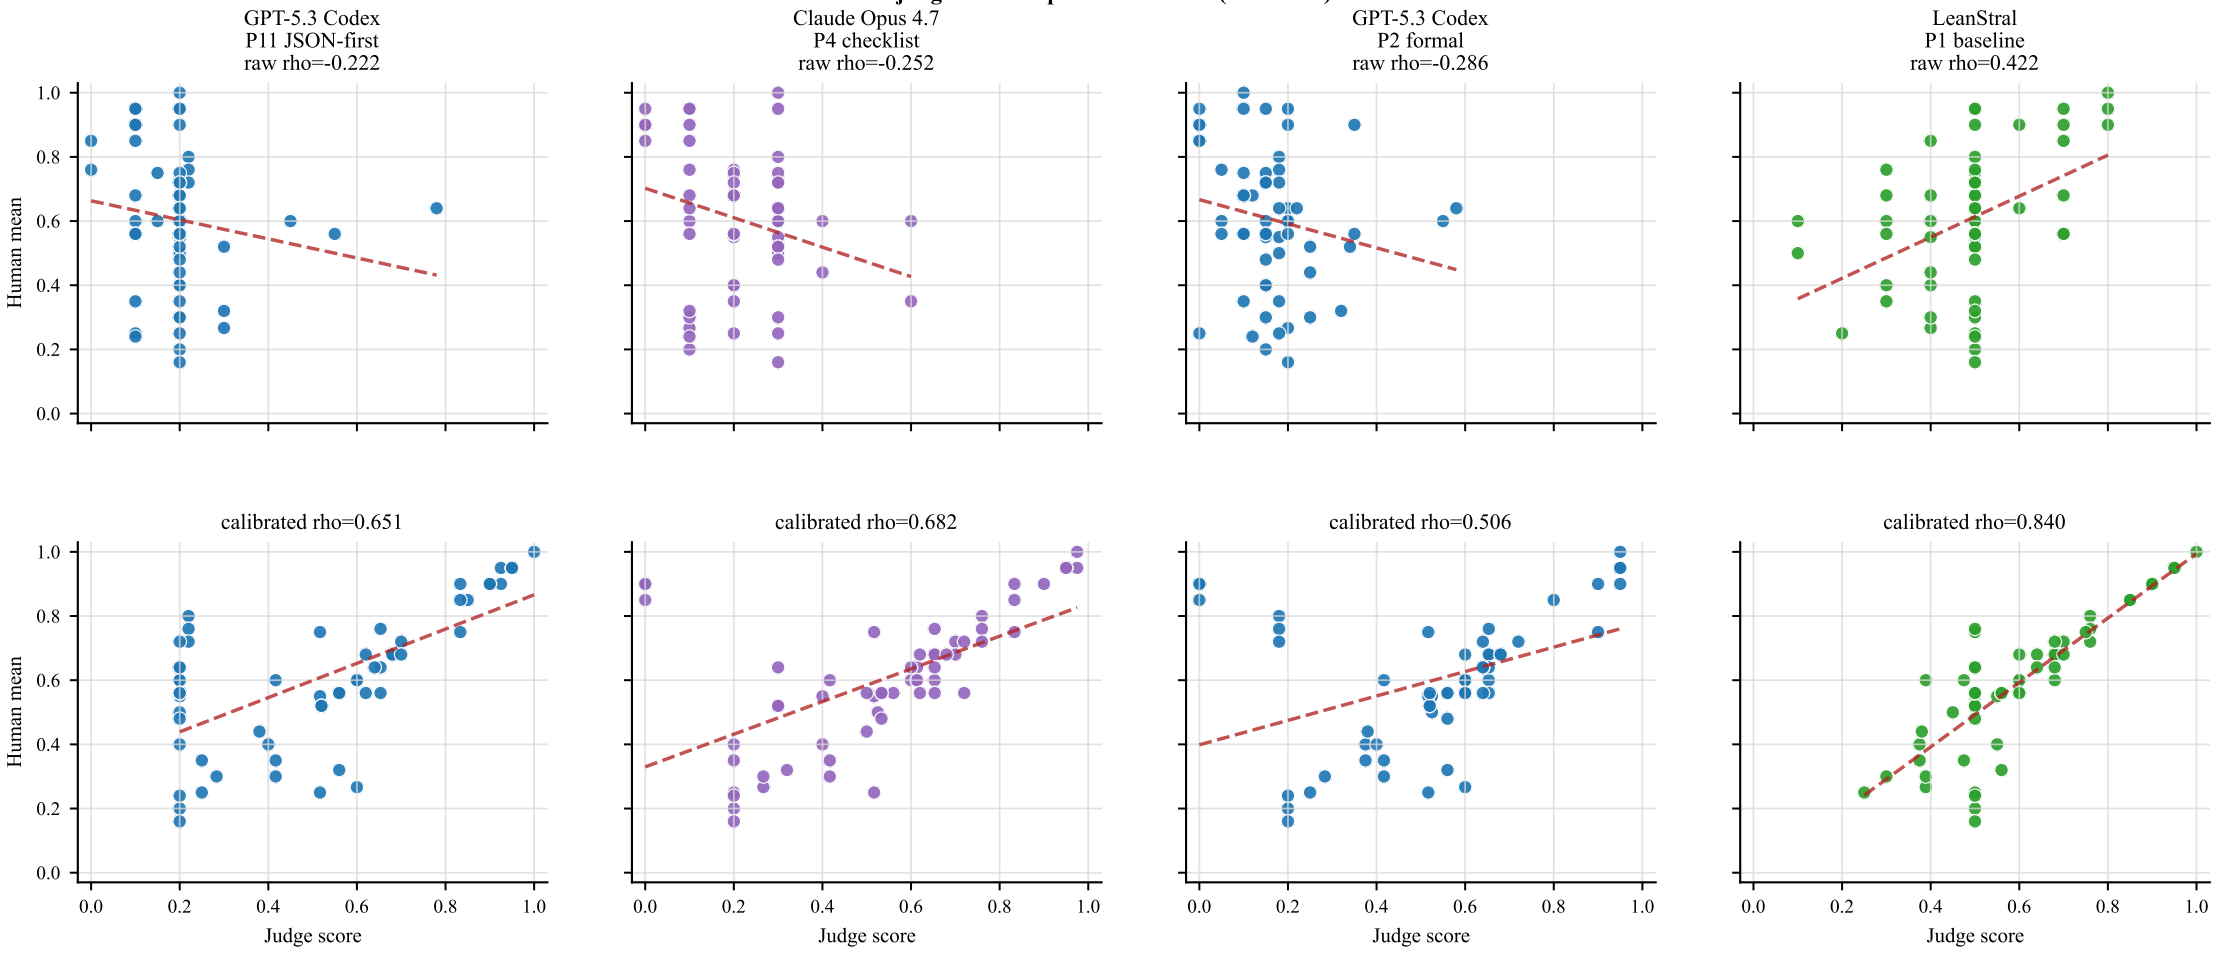
Selected judges: raw vs per-task isotonic (item-level)
GPT-5.3 Codex
P11 JSON-first
raw rho=-0.222
Claude Opus 4.7
P4 checklist
raw rho=-0.252
GPT-5.3 Codex
P2 formal
raw rho=-0.286
LeanStral
P1 baseline
raw rho=0.422
calibrated rho=0.651
calibrated rho=0.682
calibrated rho=0.506
calibrated rho=0.840
Human mean
Judge score
0.0
0.2
0.4
0.6
0.8
1.0