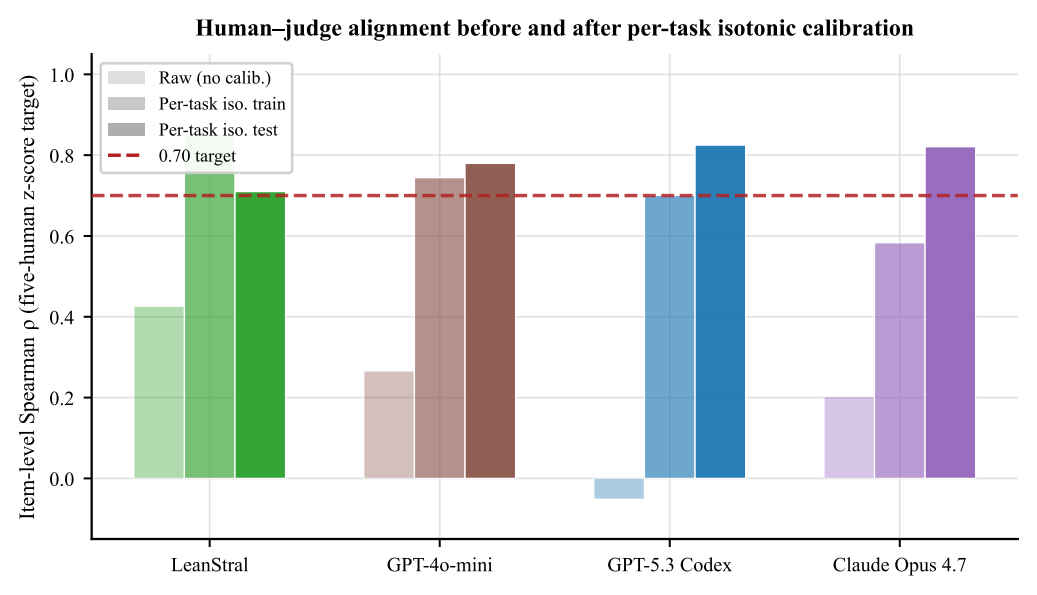
Human–judge alignment before and after per-task isotonic calibration
Item-level Spearman ρ (five-human z-score target)
1.0
0.8
0.6
0.4
0.2
0.0
Raw (no calib.)
Per-task iso. train
Per-task iso. test
0.70 target
LeanStral
GPT-4o-mini
GPT-5.3 Codex
Claude Opus 4.7