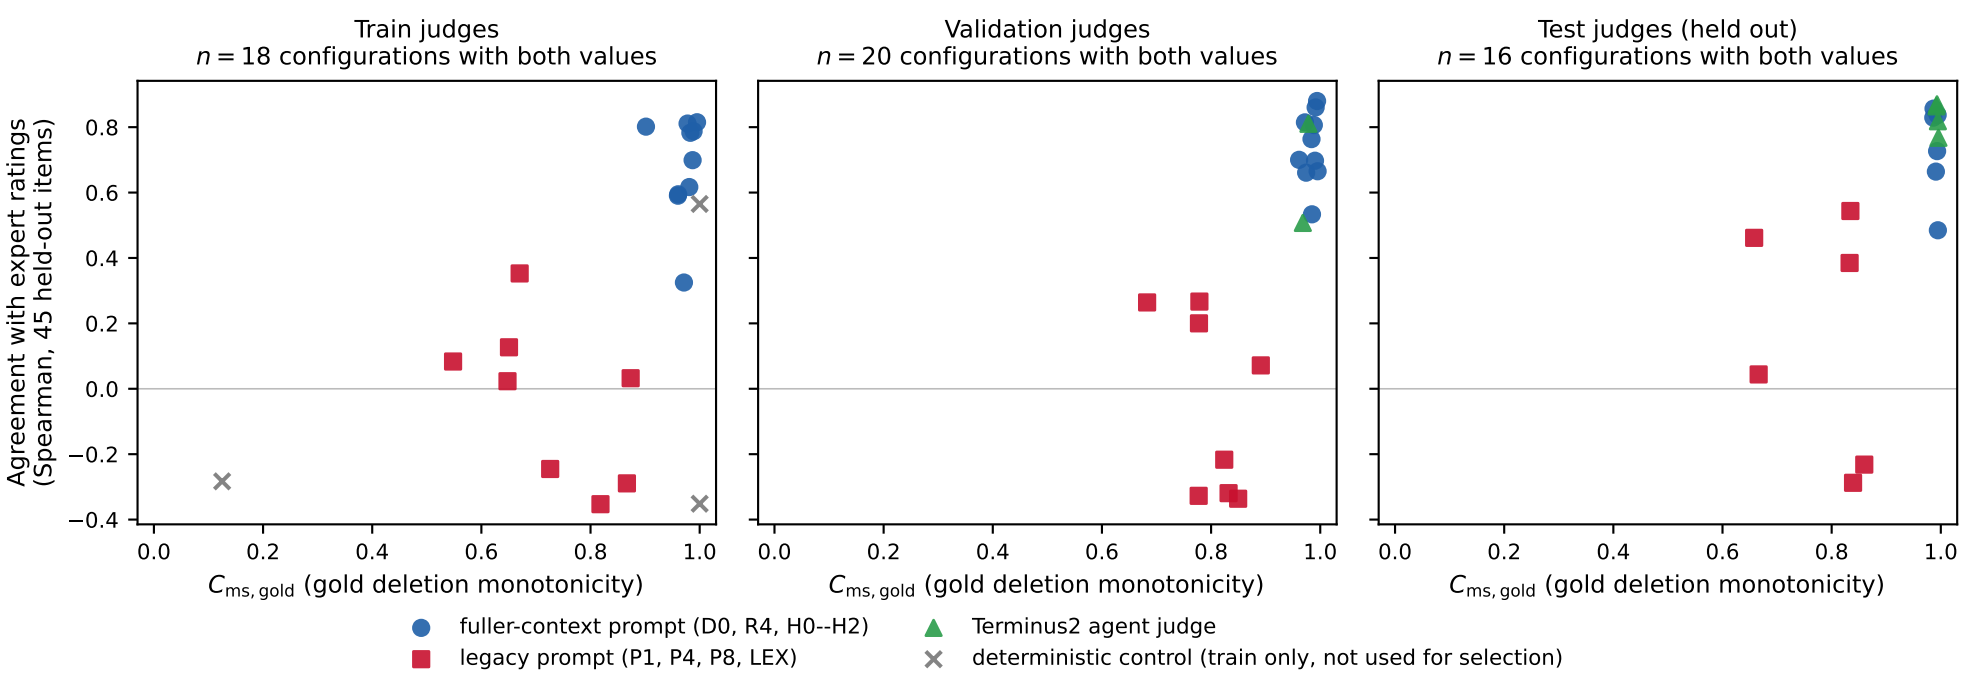
Train judges
n = 18 configurations with both values
Validation judges
n = 20 configurations with both values
Test judges (held out)
n = 16 configurations with both values
Agreement with expert ratings
(Spearman, 45 held-out items)
0.8
0.6
0.4
0.2
0.0
−0.2
−0.4
0.0
0.2
0.4
0.6
0.8
1.0
C_ms, gold (gold deletion monotonicity)
fuller-context prompt (D0, R4, H0--H2)
legacy prompt (P1, P4, P8, LEX)
Terminus2 agent judge
deterministic control (train only, not used for selection)